



Credit Card Fraud Detection

Using the Machine Learning Classification Algorithms to detect Credit Card Fraudulent Activities

Credit Card Fraud Detection

Dataset :-- <https://drive.google.com/drive/folders/14hpoXhDQgP2x5gGAcYwunA6tMlNPZEKG>

Most of the approaches involve building model on such imbalanced data, and thus fails to produce results on real-time new data because of overfitting on training data and a bias towards the majoritarian class of legitimate transactions. Thus, we can see this as an anomaly detection problem.

- 1) What time does the Credit Card Frauds usually take place?
- 2) What are the general trends of amounts for Credit Card Fraud Transactions?
- 3) How do we balance the data to not let the model overfit on legitimate transactions?

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import SGDClassifier

from mlxtend.plotting import plot_learning_curves
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.metrics import precision_score, recall_score, f1_score,
roc_auc_score, accuracy_score, classification_report
from sklearn.model_selection import KFold, StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer, matthews_corrcoef
from sklearn.feature_selection import mutual_info_classif

import warnings
warnings.filterwarnings("ignore")
```

Data Understanding

```
# read the dataset using pandas library
credit_card_data = pd.read_csv(r"C:\Users\user\Documents\
creditcard.csv")
```

credit_card_data

| | Time | V1 | V2 | V3 | V4 | |
|--------|----------|------------|-----------|-----------|-----------|-----------|
| V5 \ | | | | | | |
| 0 | 0.0 | -1.359807 | -0.072781 | 2.536347 | 1.378155 | -0.338321 |
| 1 | 0.0 | 1.191857 | 0.266151 | 0.166480 | 0.448154 | 0.060018 |
| 2 | 1.0 | -1.358354 | -1.340163 | 1.773209 | 0.379780 | -0.503198 |
| 3 | 1.0 | -0.966272 | -0.185226 | 1.792993 | -0.863291 | -0.010309 |
| 4 | 2.0 | -1.158233 | 0.877737 | 1.548718 | 0.403034 | -0.407193 |
| ... | ... | ... | ... | ... | ... | ... |
| 284802 | 172786.0 | -11.881118 | 10.071785 | -9.834783 | -2.066656 | -5.364473 |
| 284803 | 172787.0 | -0.732789 | -0.055080 | 2.035030 | -0.738589 | 0.868229 |
| 284804 | 172788.0 | 1.919565 | -0.301254 | -3.249640 | -0.557828 | 2.630515 |
| 284805 | 172788.0 | -0.240440 | 0.530483 | 0.702510 | 0.689799 | -0.377961 |
| 284806 | 172792.0 | -0.533413 | -0.189733 | 0.703337 | -0.506271 | -0.012546 |

| | V6 | V7 | V8 | V9 | ... | V21 |
|----------|-----------|-----------|-----------|-----------|-----|-----------|
| V22 \ | | | | | | |
| 0 | 0.462388 | 0.239599 | 0.098698 | 0.363787 | ... | -0.018307 |
| 0.277838 | | | | | | |
| 1 | -0.082361 | -0.078803 | 0.085102 | -0.255425 | ... | -0.225775 |
| 0.638672 | | | | | | |
| 2 | 1.800499 | 0.791461 | 0.247676 | -1.514654 | ... | 0.247998 |
| 0.771679 | | | | | | |
| 3 | 1.247203 | 0.237609 | 0.377436 | -1.387024 | ... | -0.108300 |
| 0.005274 | | | | | | |
| 4 | 0.095921 | 0.592941 | -0.270533 | 0.817739 | ... | -0.009431 |
| 0.798278 | | | | | | |
| ... | ... | ... | ... | ... | ... | ... |
| . | | | | | | |
| 284802 | -2.606837 | -4.918215 | 7.305334 | 1.914428 | ... | 0.213454 |
| 0.111864 | | | | | | |
| 284803 | 1.058415 | 0.024330 | 0.294869 | 0.584800 | ... | 0.214205 |
| 0.924384 | | | | | | |
| 284804 | 3.031260 | -0.296827 | 0.708417 | 0.432454 | ... | 0.232045 |
| 0.578229 | | | | | | |
| 284805 | 0.623708 | -0.686180 | 0.679145 | 0.392087 | ... | 0.265245 |
| 0.800049 | | | | | | |
| 284806 | -0.649617 | 1.577006 | -0.414650 | 0.486180 | ... | 0.261057 |
| 0.643078 | | | | | | |

| | V23 | V24 | V25 | V26 | V27 | V28 |
|----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Amount \ | | | | | | |
| 0 | -0.110474 | 0.066928 | 0.128539 | -0.189115 | 0.133558 | -0.021053 |
| 149.62 | | | | | | |
| 1 | 0.101288 | -0.339846 | 0.167170 | 0.125895 | -0.008983 | 0.014724 |
| 2.69 | | | | | | |
| 2 | 0.909412 | -0.689281 | -0.327642 | -0.139097 | -0.055353 | -0.059752 |
| 378.66 | | | | | | |
| 3 | -0.190321 | -1.175575 | 0.647376 | -0.221929 | 0.062723 | 0.061458 |
| 123.50 | | | | | | |
| 4 | -0.137458 | 0.141267 | -0.206010 | 0.502292 | 0.219422 | 0.215153 |
| 69.99 | | | | | | |
| ... | ... | ... | ... | ... | ... | ... |
| ... | | | | | | |
| 284802 | 1.014480 | -0.509348 | 1.436807 | 0.250034 | 0.943651 | 0.823731 |
| 0.77 | | | | | | |
| 284803 | 0.012463 | -1.016226 | -0.606624 | -0.395255 | 0.068472 | -0.053527 |
| 24.79 | | | | | | |
| 284804 | -0.037501 | 0.640134 | 0.265745 | -0.087371 | 0.004455 | -0.026561 |
| 67.88 | | | | | | |
| 284805 | -0.163298 | 0.123205 | -0.569159 | 0.546668 | 0.108821 | 0.104533 |
| 10.00 | | | | | | |
| 284806 | 0.376777 | 0.008797 | -0.473649 | -0.818267 | -0.002415 | 0.013649 |
| 217.00 | | | | | | |

| | Class |
|--------|-------|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 0 |
| ... | ... |
| 284802 | 0 |
| 284803 | 0 |
| 284804 | 0 |
| 284805 | 0 |
| 284806 | 0 |

[284807 rows x 31 columns]

getting the first 5 row's values
credit_card_data.head()

| | Time | V1 | V2 | V3 | V4 | V5 | V6 |
|----------|------|-----------|-----------|----------|----------|-----------|-----------|
| V7 \ | | | | | | | |
| 0 | 0.0 | -1.359807 | -0.072781 | 2.536347 | 1.378155 | -0.338321 | 0.462388 |
| 0.239599 | | | | | | | |
| 1 | 0.0 | 1.191857 | 0.266151 | 0.166480 | 0.448154 | 0.060018 | -0.082361 |
| 0.078803 | | | | | | | |

```

2  1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499
0.791461
3  1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203
0.237609
4  2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921
0.592941

```

```

      V8      V9  ...      V21      V22      V23      V24
V25 \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928
0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846
0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -
0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575
0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -
0.206010

```

```

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053  149.62      0
1  0.125895 -0.008983  0.014724   2.69      0
2 -0.139097 -0.055353 -0.059752  378.66      0
3 -0.221929  0.062723  0.061458  123.50      0
4  0.502292  0.219422  0.215153   69.99      0

```

[5 rows x 31 columns]

getting the last 5 row's values

```
credit_card_data.tail()
```

```

      Time      V1      V2      V3      V4
V5 \
284802  172786.0 -11.881118  10.071785 -9.834783 -2.066656 -5.364473
284803  172787.0  -0.732789  -0.055080  2.035030 -0.738589  0.868229
284804  172788.0  1.919565  -0.301254 -3.249640 -0.557828  2.630515
284805  172788.0  -0.240440  0.530483  0.702510  0.689799 -0.377961
284806  172792.0  -0.533413  -0.189733  0.703337 -0.506271 -0.012546

```

```

      V6      V7      V8      V9  ...      V21
V22 \
284802 -2.606837 -4.918215  7.305334  1.914428  ...  0.213454
0.111864
284803  1.058415  0.024330  0.294869  0.584800  ...  0.214205
0.924384

```

```

284804  3.031260 -0.296827  0.708417  0.432454  ...  0.232045
0.578229
284805  0.623708 -0.686180  0.679145  0.392087  ...  0.265245
0.800049
284806 -0.649617  1.577006 -0.414650  0.486180  ...  0.261057
0.643078

```

| | V23 | V24 | V25 | V26 | V27 | V28 |
|----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Amount \ | | | | | | |
| 284802 | 1.014480 | -0.509348 | 1.436807 | 0.250034 | 0.943651 | 0.823731 |
| 0.77 | | | | | | |
| 284803 | 0.012463 | -1.016226 | -0.606624 | -0.395255 | 0.068472 | -0.053527 |
| 24.79 | | | | | | |
| 284804 | -0.037501 | 0.640134 | 0.265745 | -0.087371 | 0.004455 | -0.026561 |
| 67.88 | | | | | | |
| 284805 | -0.163298 | 0.123205 | -0.569159 | 0.546668 | 0.108821 | 0.104533 |
| 10.00 | | | | | | |
| 284806 | 0.376777 | 0.008797 | -0.473649 | -0.818267 | -0.002415 | 0.013649 |
| 217.00 | | | | | | |

| | Class |
|--------|-------|
| 284802 | 0 |
| 284803 | 0 |
| 284804 | 0 |
| 284805 | 0 |
| 284806 | 0 |

```
[5 rows x 31 columns]
```

```

#getting the information about the whole dataset
credit_card_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null  float64
 1   V1          284807 non-null  float64
 2   V2          284807 non-null  float64
 3   V3          284807 non-null  float64
 4   V4          284807 non-null  float64
 5   V5          284807 non-null  float64
 6   V6          284807 non-null  float64
 7   V7          284807 non-null  float64
 8   V8          284807 non-null  float64
 9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64

```

```
13 V13      284807 non-null float64
14 V14      284807 non-null float64
15 V15      284807 non-null float64
16 V16      284807 non-null float64
17 V17      284807 non-null float64
18 V18      284807 non-null float64
19 V19      284807 non-null float64
20 V20      284807 non-null float64
21 V21      284807 non-null float64
22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class    284807 non-null int64
```

```
dtypes: float64(30), int64(1)
```

```
memory usage: 67.4 MB
```

#to check the null value in each column

```
credit_card_data.isnull().sum()
```

```
Time      0
V1         0
V2         0
V3         0
V4         0
V5         0
V6         0
V7         0
V8         0
V9         0
V10        0
V11        0
V12        0
V13        0
V14        0
V15        0
V16        0
V17        0
V18        0
V19        0
V20        0
V21        0
V22        0
V23        0
V24        0
V25        0
```

```
V26      0
V27      0
V28      0
Amount   0
Class    0
dtype: int64
```

Data Preparation

The Data does not have any missing values and hence, need not be handled.

The Data has only Target Variable Class as the categorical variable.

Remaining Features are numerical and need to be only standardized for comparison after balancing the dataset

```
credit_card_data.describe()
```

| | Time | V1 | V2 | V3 |
|-------|---------------|---------------|---------------|---------------|
| V4 \ | | | | |
| count | 284807.000000 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 |
| mean | 94813.859575 | 3.918649e-15 | 5.682686e-16 | -8.761736e-15 |
| std | 47488.145955 | 1.958696e+00 | 1.651309e+00 | 1.516255e+00 |
| min | 0.000000 | -5.640751e+01 | -7.271573e+01 | -4.832559e+01 |
| 25% | 54201.500000 | -9.203734e-01 | -5.985499e-01 | -8.903648e-01 |
| 50% | 84692.000000 | 1.810880e-02 | 6.548556e-02 | 1.798463e-01 |
| 75% | 139320.500000 | 1.315642e+00 | 8.037239e-01 | 1.027196e+00 |
| max | 172792.000000 | 2.454930e+00 | 2.205773e+01 | 9.382558e+00 |

| | V5 | V6 | V7 | V8 |
|-------|---------------|---------------|---------------|---------------|
| V9 \ | | | | |
| count | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 |
| mean | -1.552103e-15 | 2.040130e-15 | -1.698953e-15 | -1.893285e-16 |
| std | 1.380247e+00 | 1.332271e+00 | 1.237094e+00 | 1.194353e+00 |
| min | -1.137433e+02 | -2.616051e+01 | -4.355724e+01 | -7.321672e+01 |
| 25% | -6.915971e-01 | -7.682956e-01 | -5.540759e-01 | -2.086297e-01 |


```

6.430976e-01
50%    -5.433583e-02 -2.741871e-01  4.010308e-02  2.235804e-02 -
5.142873e-02
75%    6.119264e-01  3.985649e-01  5.704361e-01  3.273459e-01
5.971390e-01
max     3.480167e+01  7.330163e+01  1.205895e+02  2.000721e+01
1.559499e+01

```

| | ... | V21 | V22 | V23 | V24 | \ |
|-------|-----|---------------|---------------|---------------|---------------|---|
| count | ... | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 | |
| mean | ... | 1.473120e-16 | 8.042109e-16 | 5.282512e-16 | 4.456271e-15 | |
| std | ... | 7.345240e-01 | 7.257016e-01 | 6.244603e-01 | 6.056471e-01 | |
| min | ... | -3.483038e+01 | -1.093314e+01 | -4.480774e+01 | -2.836627e+00 | |
| 25% | ... | -2.283949e-01 | -5.423504e-01 | -1.618463e-01 | -3.545861e-01 | |
| 50% | ... | -2.945017e-02 | 6.781943e-03 | -1.119293e-02 | 4.097606e-02 | |
| 75% | ... | 1.863772e-01 | 5.285536e-01 | 1.476421e-01 | 4.395266e-01 | |
| max | ... | 2.720284e+01 | 1.050309e+01 | 2.252841e+01 | 4.584549e+00 | |

| | V25 | V26 | V27 | V28 |
|---------------|---------------|---------------|---------------|---------------|
| Amount \ | | | | |
| count | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 | 2.848070e+05 |
| 284807.000000 | | | | |
| mean | 1.426896e-15 | 1.701640e-15 | -3.662252e-16 | -1.217809e-16 |
| 88.349619 | | | | |
| std | 5.212781e-01 | 4.822270e-01 | 4.036325e-01 | 3.300833e-01 |
| 250.120109 | | | | |
| min | -1.029540e+01 | -2.604551e+00 | -2.256568e+01 | -1.543008e+01 |
| 0.000000 | | | | |
| 25% | -3.171451e-01 | -3.269839e-01 | -7.083953e-02 | -5.295979e-02 |
| 5.600000 | | | | |
| 50% | 1.659350e-02 | -5.213911e-02 | 1.342146e-03 | 1.124383e-02 |
| 22.000000 | | | | |
| 75% | 3.507156e-01 | 2.409522e-01 | 9.104512e-02 | 7.827995e-02 |
| 77.165000 | | | | |
| max | 7.519589e+00 | 3.517346e+00 | 3.161220e+01 | 3.384781e+01 |
| 25691.160000 | | | | |

| | Class |
|-------|---------------|
| count | 284807.000000 |
| mean | 0.001727 |
| std | 0.041527 |
| min | 0.000000 |
| 25% | 0.000000 |
| 50% | 0.000000 |
| 75% | 0.000000 |
| max | 1.000000 |

[8 rows x 31 columns]

```
# checking how many transactions are fraudulent and legitimate
credit_card_data['Class'].value_counts()
```

```
0    284315
1       492
Name: Class, dtype: int64
```

```
#seperating the values of fraud and legit data
legit = credit_card_data[credit_card_data.Class == 0]
fraud = credit_card_data[credit_card_data.Class == 1]
```

```
print(legit.shape)
print(fraud.shape)
```

```
(284315, 31)
(492, 31)
```

```
legit.Amount.describe()
```


```
count    284315.000000
mean       88.291022
std       250.105092
min         0.000000
25%         5.650000
50%        22.000000
75%        77.050000
max       25691.160000
Name: Amount, dtype: float64
```

```
fraud.Amount.describe()
```

```
count      492.000000
mean       122.211321
std        256.683288
min         0.000000
25%         1.000000
50%         9.250000
75%        105.890000
max        2125.870000
Name: Amount, dtype: float64
```

```
def countplot_data(data, feature):
    plt.figure(figsize = (10, 10))
    sns.countplot(x = feature, data = data)
    plt.show
```


```
countplot_data(credit_card_data, credit_card_data.Class)
```




```
def pairplot_data_grid(data, feature1, feature2, target):  
    sns.FacetGrid(data, hue=target, size=6).map(plt.scatter, feature1,  
    feature2).add_legend()  
    plt.show()
```

Relationship of fraud transactions with amount of money

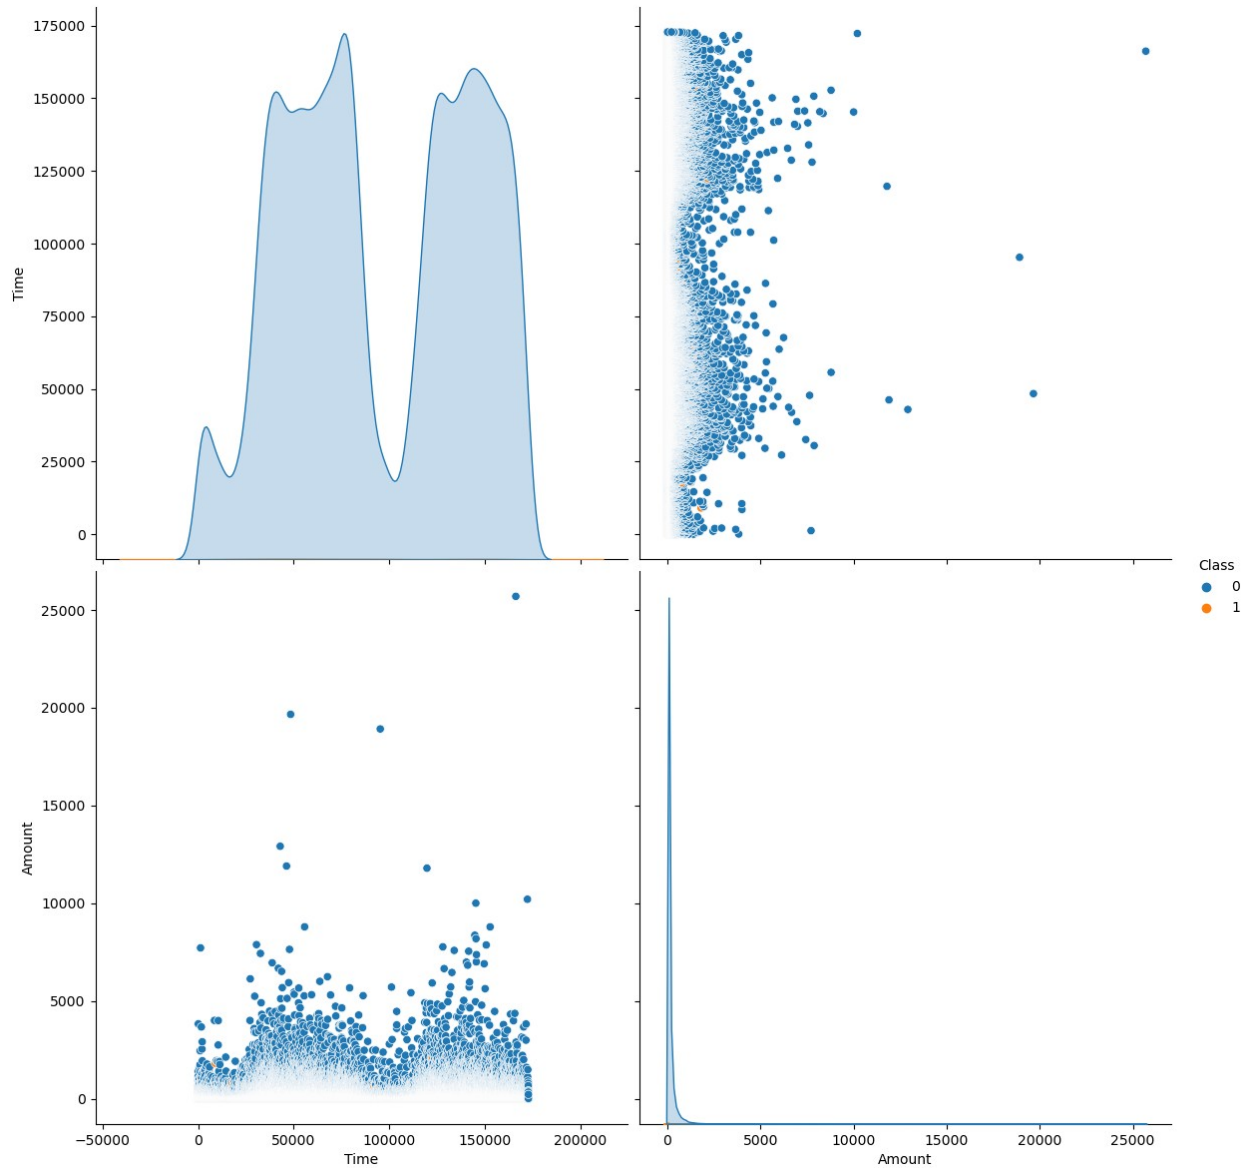
```
pairplot_data_grid(credit_card_data, "Time", "Amount", "Class")
```



```
pairplot_data_grid(credit_card_data, "Amount", "Time", "Class")
```




```
credit_card_data_refine = credit_card_data[["Time", "Amount",  
"Class"]]  
sns.pairplot(credit_card_data_refine, hue="Class", size=6)  
plt.show()
```



The relationship between Time and Transaction.

```
sns.FacetGrid(credit_card_data_refine, hue="Class",
size=6).map(sns.distplot, "Class").add_legend()
plt.show()
```



Modelling

Study the Feature Correlations of the given data

Plot a Heatmap


Run GridSearch on the Data

Fine Tune the Classifiers

Create Pipelines for evaluation

```
plt.figure(figsize = (20, 20))  
credit_card_data_corr = credit_card_data.corr()  
sns.heatmap(credit_card_data_corr)
```

```
<AxesSubplot:>
```



Now splitting the dataset into Features and Targets

```
X = credit_card_data.drop(labels='Class', axis=1) # Features
Y = credit_card_data.loc[:, 'Class']           # Target Variable
```

Now splitting the dataset into Train data and Test data


```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.3, random_state=1, stratify=Y)
```

Balancing the fraud and legitimate transactions in data

```
# Use Synthetic Minority Oversampling Techniques.
sm = SMOTE(random_state = 42)
X_res, Y_res = sm.fit_resample(X_train, Y_train)

sns.countplot(Y_res)

<AxesSubplot:xlabel='Class', ylabel='count'>
```



Evaluation

We make use of AUC-ROC Score, Classification Report, Accuracy and F1-Score to evaluate the performance of the classifiers

```
# Evaluation of Classifiers
def grid_eval(grid_clf):
    """
        Method to Compute the best score and parameters computed by
        grid search
        Parameter:
            grid_clf: The Grid Search Classifier
    """
    print("Best Score", grid_clf.best_score_)
    print("Best Parameter", grid_clf.best_params_)

def evaluation(Y_test, grid_clf, X_test):
    """
        Method to compute the following:
        1. Classification Report
        2. F1-score
        3. AUC-ROC score
    """
```

4. Accuracy

Parameters:

y_test: The target variable test set

grid_clf: Grid classifier selected

X_test: Input Feature Test Set

"""

```
Y_pred = grid_clf.predict(X_test)
```

```
print('CLASSIFICATION REPORT')
```

```
print(classification_report(Y_test, Y_pred))
```

```
print('AUC-ROC')
```

```
print(roc_auc_score(Y_test, Y_pred))
```

```
print('F1-Score')
```

```
print(f1_score(Y_test, Y_pred))
```

```
print('Accuracy')
```

```
print(accuracy_score(Y_test, Y_pred))
```

The parameters of each classifier are different

Hence, we do not make use of a single method and this is not to violate DRY Principles

We set pipelines for each classifier unique with parameters

```
param_grid_sgd = [{
    'model__loss': ['log'],
    'model__penalty': ['l1', 'l2'],
    'model__alpha': np.logspace(start=-3, stop=3, num=20)
}, {
    'model__loss': ['hinge'],
    'model__alpha': np.logspace(start=-3, stop=3, num=20),
    'model__class_weight': [None, 'balanced']
}]
```

```
pipeline_sgd = Pipeline([
    ('scaler', StandardScaler(copy=False)),
    ('model', SGDClassifier(max_iter=1000, tol=1e-3, random_state=1,
warm_start=True))
])
```

```
MCC_scorer = make_scorer(matthews_corrcoef)
```

```
grid_sgd = GridSearchCV(estimator=pipeline_sgd,
```

```
param_grid=param_grid_sgd,
```

```
scoring=MCC_scorer, n_jobs=-1,
```

```
pre_dispatch='2*n_jobs', cv=5, verbose=1, return_train_score=False)
```

```
grid_sgd.fit(X_res, Y_res)
```

Fitting 5 folds for each of 80 candidates, totalling 400 fits

```

GridSearchCV(cv=5,
             estimator=Pipeline(steps=[('scaler',
                                         StandardScaler(copy=False)),
                                         ('model',
                                          SGDClassifier(random_state=1,
                                                         warm_start=True))])),
             n_jobs=-1,
             param_grid=[{'model__alpha': array([1.00000000e-03,
2.06913808e-03, 4.28133240e-03, 8.85866790e-03,
1.83298071e-02, 3.79269019e-02, 7.84759970e-02, 1.62377674e-01,
3.35981829e-01, 6.95192796e-01, 1.43844989e+00,...
1.83298071e-02, 3.79269019e-02, 7.84759970e-02, 1.62377674e-01,
3.35981829e-01, 6.95192796e-01, 1.43844989e+00, 2.97635144e+00,
6.15848211e+00, 1.27427499e+01, 2.63665090e+01, 5.45559478e+01,
1.12883789e+02, 2.33572147e+02, 4.83293024e+02,
1.00000000e+03])},
                        {'model__class_weight': [None, 'balanced'],
                         'model__loss': ['hinge']}]},
             scoring=make_scorer(matthews_corrcoef), verbose=1)

```

```
grid_eval(grid_sgd)
```

Best Score 0.9560162686072134

Best Parameter {'model__alpha': 0.001, 'model__loss': 'log',
'model__penalty': 'l1'}

```
evaluation(Y_test, grid_sgd, X_test)
```

CLASSIFICATION REPORT

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.99 | 1.00 | 85295 |
| 1 | 0.14 | 0.91 | 0.25 | 148 |
| accuracy | | | 0.99 | 85443 |
| macro avg | 0.57 | 0.95 | 0.62 | 85443 |
| weighted avg | 1.00 | 0.99 | 0.99 | 85443 |

AUC-ROC

0.9479720619851928

F1-Score

0.2460973370064279

Accuracy

0.990391254988706

```

pipeline_rf = Pipeline([
    ('model', RandomForestClassifier(n_jobs=-1, random_state=1))]
)
param_grid_rf = {'model__n_estimators': [75]}
grid_rf = GridSearchCV(estimator=pipeline_rf,
                       param_grid=param_grid_rf,

```

```

        scoring=MCC_scorer, n_jobs=-1,
pre_dispatch='2*n_jobs', cv=5, verbose=1, return_train_score=False)
grid_rf.fit(X_res, Y_res)

```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```

GridSearchCV(cv=5,
             estimator=Pipeline(steps=[('model',
RandomForestClassifier(n_jobs=-1,
random_state=1))]),
             n_jobs=-1, param_grid={'model__n_estimators': [75]},
             scoring=make_scorer(matthews_corrcoef), verbose=1)

```

```
grid_eval(grid_rf)
```

Best Score 0.9997538267139271

Best Parameter {'model__n_estimators': 75}

```
evaluation(Y_test, grid_rf, X_test)
```

CLASSIFICATION REPORT

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 85295 |
| 1 | 0.90 | 0.86 | 0.88 | 148 |
| accuracy | | | 1.00 | 85443 |
| macro avg | 0.95 | 0.93 | 0.94 | 85443 |
| weighted avg | 1.00 | 1.00 | 1.00 | 85443 |

AUC-ROC

0.9323445023075716

F1-Score

0.879725085910653

Accuracy

0.9995903701883126

```

pipeline_lr = Pipeline([
    ('model', LogisticRegression(random_state=1))
])
param_grid_lr = {'model__penalty': ['l2'],
                 'model__class_weight': [None, 'balanced']}
grid_lr = GridSearchCV(estimator=pipeline_lr,
                       param_grid=param_grid_lr,
                       scoring=MCC_scorer, n_jobs=-1,
pre_dispatch='2*n_jobs', cv=5, verbose=1, return_train_score=False)
grid_lr.fit(X_res, Y_res)

```

Fitting 5 folds for each of 2 candidates, totalling 10 fits

```

GridSearchCV(cv=5,
             estimator=Pipeline(steps=[('model',
LogisticRegression(random_state=1))]),
             n_jobs=-1,
             param_grid={'model__class_weight': [None, 'balanced'],
                         'model__penalty': ['l2']},
             scoring=make_scorer(matthews_corrcoef), verbose=1)

grid_eval(grid_lr)

Best Score 0.959816277887179
Best Parameter {'model__class_weight': None, 'model__penalty': 'l2'}

evaluation(Y_test, grid_lr, X_test)

CLASSIFICATION REPORT

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.99 | 1.00 | 85295 |
| 1 | 0.15 | 0.91 | 0.26 | 148 |
| accuracy | | | 0.99 | 85443 |
| macro avg | 0.57 | 0.95 | 0.63 | 85443 |
| weighted avg | 1.00 | 0.99 | 0.99 | 85443 |

```

AUC-ROC
0.948212404326479
F1-Score
0.2557251908396946
Accuracy
0.9908711070538253

```

Conclusion

The K-Nearest Neighbors Classifier tuned with Grid Search with the best parameter and its counterparts to give a test accuracy of nearly 99.9% and a perfect F1-Score with minimal overfitting